

Stat 897 Fall 2017 Data Analysis Assignment 9

Penn State

Due October 29, 2017

In this assignment we will use the Boston data found in the MASS library.

1. Fit classification models in order to predict whether a given suburb has a crime rate above or below the median. Explore logistic regression, LDA, and KNN models using various subsets of the predictors. At this point in the class, you should feel fairly comfortable making such an open-ended exploration.

Describe your findings, show appropriate results, and determine why some techniques perform better or worse and which had the best performance.

```
library(MASS)
library(class)
library(glmnet)
```

```
## Loading required package: Matrix
## Loading required package: foreach
## Loaded glmnet 2.0-12
```

```
library(leaps)
library(caret)
```

```
## Loading required package: lattice
## Loading required package: ggplot2
```

```
data("Boston")
attach(Boston)
```

Let's start with parameter selection for the Boston data set. We will use forward selection, lasso and ridge here:

```
x = model.matrix(crim ~ . - 1, data = Boston)
y = Boston$crim

n = nrow(Boston)
p = ncol(Boston) - 1
set.seed(801)
trainingRows=sample(nrow(Boston), n*0.7, replace = FALSE)
train = Boston[trainingRows,]
test = Boston[-trainingRows,]
train.mat <- model.matrix(crim~ ., data = train)
test.mat <- model.matrix(crim~ ., data = test)

# Forward Selection / BIC
regfit.fwd=regsubsets(crim~.,data=train, nvmax =14, method='forward')
reg.summary = summary(regfit.fwd)
reg.summary
```

```
## Subset selection object
## Call: regsubsets.formula(crim ~ ., data = train, nvmax = 14, method = "forward")
## 13 Variables (and intercept)
##           Forced in Forced out
## zn          FALSE      FALSE
## indus        FALSE      FALSE
## chas         FALSE      FALSE
## nox          FALSE      FALSE
## rm           FALSE      FALSE
## age          FALSE      FALSE
## dis          FALSE      FALSE
## rad          FALSE      FALSE
## tax          FALSE      FALSE
## ptratio      FALSE      FALSE
## black        FALSE      FALSE
## lstat        FALSE      FALSE
## medv         FALSE      FALSE
```

```
## 1 subsets of each size up to 13
```

```
## Selection Algorithm: forward
```


```
##           zn indus chas nox rm age dis rad tax ptratio black lstat medv
## 1 ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 2 ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 3 ( 1 ) " " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 4 ( 1 ) "*" " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 5 ( 1 ) "*" " " " " " " " " " " " " " " " " " " " " " " " " " " "
## 6 ( 1 ) "*" " " " " " " "*" " " " "*" " " " " " " " " " " " " " " "
## 7 ( 1 ) "*" "*" " " " " "*" " " " "*" " " " " " " " " " " " " " " "
## 8 ( 1 ) "*" "*" " " " "*" "*" " " " "*" " " " " " " " " " " " " " "
## 9 ( 1 ) "*" "*" "*" "*" "*" " " " "*" " " " " " " " " " " " " " "
## 10 ( 1 ) "*" "*" "*" "*" "*" " " " "*" " " " " "*" " " " " " " " "
## 11 ( 1 ) "*" "*" "*" "*" "*" " " " "*" " " " " "*" " " " "*" " " "
## 12 ( 1 ) "*" "*" "*" "*" "*" " " " "*" " " "*" "*" "*" " " " "*" " "
## 13 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" " " " "*" " " "
```

```
plot(reg.summary$bic, xlab = "Number of Variables", ylab = "BIC", type = 'l', main = 'Forward Step - Performance')
which.min (reg.summary$bic )
```


```
## [1] 3
```

```
points (which.min (reg.summary$bic ), reg.summary$bic[which.min (reg.summary$bic )], col = "red", cex = 2,
```

Forward Step – Performance Measure



```
#LASSO  
grid = 10^ seq (10,-2, length =100)  
cv.lasso = cv.glmnet(x, y, type.measure = "mse", nfolds=10)  
plot(cv.lasso)
```



```
bestlam.lasso=cv.lasso$lambda.min #find the best tuning parameter

fit.lasso <- glmnet(train.mat, train$crim, alpha = 1, lambda = grid, thresh = 1e-12)
pred.lasso=predict (fit.lasso, s=bestlam.lasso, newx=test.mat)
mean(( pred.lasso - test$crim)^2)
```

```
## [1] 104.9072
```

```
final.lasso=glmnet(x,y,alpha=1) #fit on the entire data set to extract coef
lasso.coef=predict(final.lasso,type="coefficients",s=bestlam.lasso)[1:14,]
lasso.coef
```

```
## (Intercept)          zn          indus          chas          nox
## 12.666630982  0.036235999 -0.070797081 -0.585263575 -6.966972654
##          rm          age          dis          rad          tax
##  0.229498615  0.000000000 -0.787816119  0.514355222  0.000000000
##      ptratio        black        lstat        medv
## -0.188016174 -0.007548647  0.125278643 -0.157926207
```

```
length(lasso.coef[lasso.coef !=0])
```

```
## [1] 12
```


```
lasso.coef[lasso.coef!=0] #contains 11 variables in our model
```

```
## (Intercept)          zn          indus          chas          nox
## 12.666630982  0.036235999 -0.070797081 -0.585263575 -6.966972654
##          rm          dis          rad      ptratio        black
##  0.229498615 -0.787816119  0.514355222 -0.188016174 -0.007548647
```

```
##          lstat          medv
## 0.125278643 -0.157926207
```

```
#Ridge regression
```

```
cv.ridge = cv.glmnet(x, y, alpha=0, type.measure = "mse", nfolds=length(y),grouped=FALSE)
plot(cv.ridge)
```



```
bestlam.ridge=cv.ridge$lambda.min #find the best tuning parameter
```

```
fit.ridge = glmnet(train.mat, train$crim, alpha = 0, lambda = grid, thresh = 1e-12)
pred.ridge = predict (fit.ridge, s=bestlam.ridge, newx=test.mat)
mean(( pred.ridge - test$crim)^2)
```

```
## [1] 106.0669
```

```
#fit on the full data
```

```
ridge.coef=predict(final.ridge,type="coefficients",s=bestlam.ridge)[1:14,]
ridge.coef
```

```
## (Intercept)          zn          indus          chas          nox
## 8.617905279 0.032352168 -0.081183885 -0.739986141 -5.095661105
##          rm          age          dis          rad          tax
## 0.328170831 0.002074971 -0.683786238 0.414237411 0.003695600
##          ptratio        black          lstat          medv
## -0.127614747 -0.008532788 0.142710654 -0.136308133
```

```
#contains all variables in our model
```

```
## (Intercept)          zn          indus          chas          nox
```

```
## 8.617905279 0.032352168 -0.081183885 -0.739986141 -5.095661105
##          rm          age          dis          rad          tax
## 0.328170831 0.002074971 -0.683786238 0.414237411 0.003695600
##      ptratio        black        lstat        medv
## -0.127614747 -0.008532788 0.142710654 -0.136308133
```

Based on the results above we have the following: Lasso: selects model with 11 variables: zn + indus + chas + nox + rm + dis + rad + ptratio + black + lstat + medv

Forward selection: selects model with 3 variables: rad + black + lstat

Ridge: selects all parameters - we will ignore this as we get smaller models with better Test MSE with Lasso.

```
# do some data processing to prepare the categorical variable
```

```
crim_median <- rep(0, length(crim))
crim_median[crim > median(crim)] <- 1
```

```
# add new column to the data frame
```

```
Boston <- data.frame(Boston, crim_median)
```

```
train <- 1:(length(crim) * 0.7)
test <- (length(train)+ 1):length(crim)
Boston.train <- Boston[train, ]
Boston.test <- Boston[test, ]
crim_median.test <- crim_median[test]
```

```
# Logistic Regression: Model with all parameters
```

```
fit.glm <- glm(crim_median ~ . - crim_median - crim, data = Boston, family = binomial, subset = train)
probs <- predict(fit.glm, Boston.test, type = "response")
pred.glm <- rep(0, length(probs))
pred.glm[probs > 0.5] <- 1
conf_matrix = table(pred.glm, crim_median.test)
conf_matrix
```

```
##          crim_median.test
## pred.glm    0    1
##           0    7    0
##           1   10  135
```

```
cm=confusionMatrix(data = pred.glm, reference = crim_median.test)
cm$byClass
```

```
##          Sensitivity          Specificity          Pos Pred Value
##          0.41176471          1.00000000          1.00000000
##          Neg Pred Value          Precision          Recall
##          0.93103448          1.00000000          0.41176471
##          F1          Prevalence          Detection Rate
##          0.58333333          0.11184211          0.04605263
## Detection Prevalence          Balanced Accuracy
##          0.04605263          0.70588235
```

```
mean(pred.glm == crim_median.test)
```

```
## [1] 0.9342105
```

We see that the results appear good. However specificity is 41%

```
# Logistic Regression: Model with parameters selected by lasso
```

```
fit.glm <- glm(crim_median ~ zn + indus + chas + nox + rm + dis + rad + ptratio + black + lstat + medv,
```

```
probs <- predict(fit.glm, Boston.test, type = "response")
pred.glm <- rep(0, length(probs))
pred.glm[probs > 0.5] <- 1
table(pred.glm, crim_median.test)
```

```
##          crim_median.test
## pred.glm    0    1
##           0    7    0
##           1   10  135
```

```
cm=confusionMatrix(data = pred.glm, reference = crim_median.test)
cm$byClass
```

```
##          Sensitivity      Specificity      Pos Pred Value
##          0.41176471      1.00000000      1.00000000
##      Neg Pred Value      Precision      Recall
##          0.93103448      1.00000000      0.41176471
##              F1      Prevalence      Detection Rate
##          0.58333333      0.11184211      0.04605263
## Detection Prevalence      Balanced Accuracy
##          0.04605263      0.70588235
```

```
mean(pred.glm == crim_median.test)
```

```
## [1] 0.9342105
```

Results similar to when we use all params.

```
# Logistic Regression: Model with parameters selected by forward selection
fit.glm <- glm(crim_median ~ rad + black + lstat, data = Boston, family = binomial, subset = train)
probs <- predict(fit.glm, Boston.test, type = "response")
pred.glm <- rep(0, length(probs))
pred.glm[probs > 0.5] <- 1
table(pred.glm, crim_median.test)
```

```
##          crim_median.test
## pred.glm    0    1
##           0   14    2
##           1    3  133
```

```
cm_lr=confusionMatrix(data = pred.glm, reference = crim_median.test)
cm_lr$byClass
```

```
##          Sensitivity      Specificity      Pos Pred Value
##          0.82352941      0.98518519      0.87500000
##      Neg Pred Value      Precision      Recall
##          0.97794118      0.87500000      0.82352941
##              F1      Prevalence      Detection Rate
##          0.84848485      0.11184211      0.09210526
## Detection Prevalence      Balanced Accuracy
##          0.10526316      0.90435730
```

```
mean(pred.glm == crim_median.test)
```

```
## [1] 0.9671053
```

Better results and model is simpler. Moving to LDA:

```
# LDA: Model with all parameters
fit.lda <- lda(crim_median ~ . - crim_median - crim, data = Boston, subset = train)
pred.lda <- predict(fit.lda, Boston.test)
table(pred.lda$class, crim_median.test)
```

```
##      crim_median.test
##      0      1
##  0    5    0
##  1   12  135
```

```
mean(pred.lda$class == crim_median.test)
```

```
## [1] 0.9210526
```

```
cm=confusionMatrix(data = pred.lda$class, reference = crim_median.test)
cm$byClass
```

```
##      Sensitivity      Specificity      Pos Pred Value
##      0.29411765      1.00000000      1.00000000
##      Neg Pred Value      Precision      Recall
##      0.91836735      1.00000000      0.29411765
##      F1      Prevalence      Detection Rate
##      0.45454545      0.11184211      0.03289474
## Detection Prevalence      Balanced Accuracy
##      0.03289474      0.64705882
```

Poor results when we use all params.

```
# LDA: Model with parameters selected by lasso
fit.lda <- lda(crim_median ~ zn + indus + chas + nox + rm + dis + rad + ptratio + black + lstat + medv,
pred.lda <- predict(fit.lda, Boston.test)
table(pred.lda$class, crim_median.test)
```

```
##      crim_median.test
##      0      1
##  0    4    0
##  1   13  135
```

```
mean(pred.lda$class == crim_median.test)
```

```
## [1] 0.9144737
```

```
cm=confusionMatrix(data = pred.lda$class, reference = crim_median.test)
cm$byClass
```

```
##      Sensitivity      Specificity      Pos Pred Value
##      0.23529412      1.00000000      1.00000000
##      Neg Pred Value      Precision      Recall
##      0.91216216      1.00000000      0.23529412
##      F1      Prevalence      Detection Rate
##      0.38095238      0.11184211      0.02631579
## Detection Prevalence      Balanced Accuracy
##      0.02631579      0.61764706
```

Poor results when we use Lasso params. Logistic Reg performed better.

```
# LDA: Model with parameters selected by forward selection
fit.lda <- lda(crim_median ~ rad + black + lstat, data = Boston, subset = train)
```



```
pred.lda <- predict(fit.lda, Boston.test)
table(pred.lda$class, crim_median.test)
```

```
##      crim_median.test
##      0      1
## 0  13      2
## 1   4    133
```

```
mean(pred.lda$class == crim_median.test)
```

```
## [1] 0.9605263
```

```
cm_lda=confusionMatrix(data = pred.lda$class, reference = crim_median.test)
cm_lda$byClass
```

```
##      Sensitivity      Specificity      Pos Pred Value
##      0.76470588      0.98518519      0.86666667
##      Neg Pred Value      Precision      Recall
##      0.97080292      0.86666667      0.76470588
##      F1      Prevalence      Detection Rate
##      0.81250000      0.11184211      0.08552632
## Detection Prevalence      Balanced Accuracy
##      0.09868421      0.87494553
```

Simplest model performed best with LDA. For both LR and LDA, simplest model performed the best. Between LDA and LR, LR has performed slightly better.

Moving to KNN.

```
# KNN: Model with all parameters
```

```
train.X <- cbind(zn, indus, chas, nox, rm, age, dis, rad, tax, ptratio, black, lstat, medv)[train, ]
test.X <- cbind(zn, indus, chas, nox, rm, age, dis, rad, tax, ptratio, black, lstat, medv)[test, ]
train.crim_median <- crim_median[train]
set.seed(1)
pred.knn <- knn(train.X, test.X, train.crim_median, k = 1)
table(pred.knn, crim_median.test)
```

```
##      crim_median.test
## pred.knn  0      1
##      0  17  93
##      1   0  42
```

```
mean(pred.knn == crim_median.test)
```

```
## [1] 0.3881579
```

```
pred.knn <- knn(train.X, test.X, train.crim_median, k = 10)
table(pred.knn, crim_median.test)
```

```
##      crim_median.test
## pred.knn  0      1
##      0  12   4
##      1   5  131
```

```
mean(pred.knn == crim_median.test)
```

```
## [1] 0.9407895
```

```
pred.knn <- knn(train.X, test.X, train.crim_median, k = 100)
table(pred.knn, crim_median.test)
```

```
##          crim_median.test
## pred.knn    0    1
##          0  15 115
##          1   2  20
mean(pred.knn == crim_median.test)
```

```
## [1] 0.2302632
```

When using all params we get the best results when $k=10$

```
# KNN: Model with parameters selected by lasso
train.X <- cbind(zn, indus, chas, nox, rm, dis, rad, ptratio, black, lstat, medv)[train, ]
test.X <- cbind(zn, indus, chas, nox, rm, dis, rad, ptratio, black, lstat, medv)[test, ]
train.crim_median <- crim_median[train]
set.seed(1)
pred.knn <- knn(train.X, test.X, train.crim_median, k = 1)
table(pred.knn, crim_median.test)
```

```
##          crim_median.test
## pred.knn    0    1
##          0   6 55
##          1  11 80
mean(pred.knn == crim_median.test)
```

```
## [1] 0.5657895
```

```
pred.knn <- knn(train.X, test.X, train.crim_median, k = 10)
table(pred.knn, crim_median.test)
```

```
##          crim_median.test
## pred.knn    0    1
##          0   8 25
##          1   9 110
mean(pred.knn == crim_median.test)
```

```
## [1] 0.7763158
```

```
pred.knn <- knn(train.X, test.X, train.crim_median, k = 100)
table(pred.knn, crim_median.test)
```

```
##          crim_median.test
## pred.knn    0    1
##          0  15 81
##          1   2 54
mean(pred.knn == crim_median.test)
```

```
## [1] 0.4539474
```

When using lasso params results are similar. The $k=10$ mean actually falls.

```
# KNN: Model with parameters selected by forward selection
train.X <- cbind(rad, black, lstat)[train, ]
test.X <- cbind(rad, black, lstat)[test, ]
train.crim_median <- crim_median[train]
set.seed(1)
```

```
pred.knn <- knn(train.X, test.X, train.crim_median, k = 1)
table(pred.knn, crim_median.test)
```

```
##          crim_median.test
## pred.knn  0  1
##          0 12 63
##          1  5 72
```

```
mean(pred.knn == crim_median.test)
```

```
## [1] 0.5526316
```

```
pred.knn <- knn(train.X, test.X, train.crim_median, k = 10)
table(pred.knn, crim_median.test)
```

```
##          crim_median.test
## pred.knn  0  1
##          0 14 50
##          1  3 85
```

```
mean(pred.knn == crim_median.test)
```

```
## [1] 0.6513158
```

```
pred.knn <- knn(train.X, test.X, train.crim_median, k = 100)
table(pred.knn, crim_median.test)
```

```
##          crim_median.test
## pred.knn  0  1
##          0 17 91
##          1  0 44
```

```
mean(pred.knn == crim_median.test)
```

```
## [1] 0.4013158
```

```
cm_lr$byClass
```

##	Sensitivity	Specificity	Pos Pred Value
##	0.82352941	0.98518519	0.87500000
##	Neg Pred Value	Precision	Recall
##	0.97794118	0.87500000	0.82352941
##	F1	Prevalence	Detection Rate
##	0.84848485	0.11184211	0.09210526
##	Detection Prevalence	Balanced Accuracy	
##	0.10526316	0.90435730	

```
cm_lda$byClass
```

##	Sensitivity	Specificity	Pos Pred Value
##	0.76470588	0.98518519	0.86666667
##	Neg Pred Value	Precision	Recall
##	0.97080292	0.86666667	0.76470588
##	F1	Prevalence	Detection Rate
##	0.81250000	0.11184211	0.08552632
##	Detection Prevalence	Balanced Accuracy	
##	0.09868421	0.87494553	

All in all it seems that the simple Logistic Regression model with parameters selected by forward selection

performs best closely followed by LDA model with parameters selected by forward selection.

2. Now repeat the exercise but classify those neighborhoods in the bottom 10% percentile with lowest crime rates. What differences do you notice between this and the previous classification task (hint: look at the confusion matrix)? Why may it be deceiving to only look at misclassification rate? What other measures can you consider?

We will setup the test and train data as described. The difference here will be that test data will only have 0 values for crime above or below median [always below]

```
test = which(Boston$crim <= quantile(Boston$crim, 0.1))
train = which(Boston$crim > quantile(Boston$crim, 0.1))

Boston.train <- Boston[train, ]
Boston.test <- Boston[test, ]
crim_median.test <- crim_median[test]

# Logistic Regression: Model with all parameters
fit.glm <- glm(crim_median ~ . - crim_median - crim, data = Boston, family = binomial, subset = train)
probs <- predict(fit.glm, Boston.test, type = "response")
pred.glm <- rep(0, length(probs))
pred.glm[probs > 0.5] <- 1
conf_matrix = table(pred.glm, crim_median.test)
conf_matrix

##          crim_median.test
## pred.glm  0
##           0 47
##           1  4

mean(pred.glm == crim_median.test)

## [1] 0.9215686
```

We can see the problem right here that - the column for true positive is missing.

We classification rate is high and doesn't tell the complete story. We get that on seeing the confusion matrix.

```
# Logistic Regression: Model with parameters selected by lasso
fit.glm <- glm(crim_median ~ zn + indus + chas + nox + rm + dis + rad + ptratio + black + lstat + medv,
Boston.test, type = "response")
pred.glm <- rep(0, length(probs))
pred.glm[probs > 0.5] <- 1
table(pred.glm, crim_median.test)

##          crim_median.test
## pred.glm  0
##           0 47
##           1  4

mean(pred.glm == crim_median.test)

## [1] 0.9215686
```

Results similar to when we use all params.

```
# Logistic Regression: Model with parameters selected by forward selection
fit.glm <- glm(crim_median ~ rad + black + lstat, data = Boston, family = binomial, subset = train)
probs <- predict(fit.glm, Boston.test, type = "response")
```

```
pred.glm <- rep(0, length(probs))
pred.glm[probs > 0.5] <- 1
table(pred.glm, crim_median.test)
```

```
##      crim_median.test
## pred.glm  0
##           0 51
```

```
mean(pred.glm == crim_median.test)
```

```
## [1] 1
```

This model classifies all test rows correctly and we see that the confusion matrix has a single entry.

LDA: Model with all parameters

```
fit.lda <- lda(crim_median ~ . - crim_median - crim, data = Boston, subset = train)
pred.lda <- predict(fit.lda, Boston.test)
table(pred.lda$class, crim_median.test)
```

```
##      crim_median.test
##      0
##  0 51
##  1  0
```

```
mean(pred.lda$class == crim_median.test)
```

```
## [1] 1
```

This model classifies all test rows correctly.

LDA: Model with parameters selected by lasso

```
fit.lda <- lda(crim_median ~ zn + indus + chas + nox + rm + dis + rad + ptratio + black + lstat + medv,
pred.lda <- predict(fit.lda, Boston.test)
table(pred.lda$class, crim_median.test)
```

```
##      crim_median.test
##      0
##  0 51
##  1  0
```

```
mean(pred.lda$class == crim_median.test)
```

```
## [1] 1
```

This model classifies all test rows correctly.

LDA: Model with parameters selected by forward selection

```
fit.lda <- lda(crim_median ~ rad + black + lstat, data = Boston, subset = train)
pred.lda <- predict(fit.lda, Boston.test)
table(pred.lda$class, crim_median.test)
```

```
##      crim_median.test
##      0
##  0 51
##  1  0
```

```
mean(pred.lda$class == crim_median.test)
```

```
## [1] 1
```

Simplest model also classified everything correctly. For LR simplest model performed the best. For LDA all performed fine.

Moving to KNN.

```
# KNN: Model with all parameters
train.X <- cbind(zn, indus, chas, nox, rm, age, dis, rad, tax, ptratio, black, lstat, medv)[train, ]
test.X <- cbind(zn, indus, chas, nox, rm, age, dis, rad, tax, ptratio, black, lstat, medv)[test, ]
train.crim_median <- crim_median[train]
set.seed(1)
pred.knn <- knn(train.X, test.X, train.crim_median, k = 1)
table(pred.knn, crim_median.test)
```

```
##          crim_median.test
## pred.knn  0
##          0 51
##          1  0
```

```
mean(pred.knn == crim_median.test)
```

```
## [1] 1
```

```
pred.knn <- knn(train.X, test.X, train.crim_median, k = 10)
table(pred.knn, crim_median.test)
```

```
##          crim_median.test
## pred.knn  0
##          0 51
##          1  0
```

```
mean(pred.knn == crim_median.test)
```

```
## [1] 1
```

```
pred.knn <- knn(train.X, test.X, train.crim_median, k = 100)
table(pred.knn, crim_median.test)
```

```
##          crim_median.test
## pred.knn  0
##          0 51
##          1  0
```

```
mean(pred.knn == crim_median.test)
```

```
## [1] 1
```

When using all params all classified correctly for all k.

```
# KNN: Model with parameters selected by lasso
train.X <- cbind(zn, indus, chas, nox, rm, dis, rad, ptratio, black, lstat, medv)[train, ]
test.X <- cbind(zn, indus, chas, nox, rm, dis, rad, ptratio, black, lstat, medv)[test, ]
train.crim_median <- crim_median[train]
set.seed(1)
pred.knn <- knn(train.X, test.X, train.crim_median, k = 1)
table(pred.knn, crim_median.test)
```

```
##          crim_median.test
## pred.knn  0
##          0 50
##          1  1
```

```

mean(pred.knn == crim_median.test)

## [1] 0.9803922

pred.knn <- knn(train.X, test.X, train.crim_median, k = 10)
table(pred.knn, crim_median.test)

##          crim_median.test
## pred.knn  0
##          0 48
##          1  3

mean(pred.knn == crim_median.test)

## [1] 0.9411765

pred.knn <- knn(train.X, test.X, train.crim_median, k = 100)
table(pred.knn, crim_median.test)

##          crim_median.test
## pred.knn  0
##          0 51
##          1  0

mean(pred.knn == crim_median.test)

## [1] 1

```

When using lasso params k=100 performed best. Other values of k classified a few incorrectly.

```

# KNN: Model with parameters selected by forward selection
train.X <- cbind(rad, black, lstat)[train, ]
test.X <- cbind(rad, black, lstat)[test, ]
train.crim_median <- crim_median[train]
set.seed(1)
pred.knn <- knn(train.X, test.X, train.crim_median, k = 1)
table(pred.knn, crim_median.test)

##          crim_median.test
## pred.knn  0
##          0 43
##          1  8

mean(pred.knn == crim_median.test)

## [1] 0.8431373

pred.knn <- knn(train.X, test.X, train.crim_median, k = 10)
table(pred.knn, crim_median.test)

##          crim_median.test
## pred.knn  0
##          0 48
##          1  3

mean(pred.knn == crim_median.test)

## [1] 0.9411765

```

```
pred.knn <- knn(train.X, test.X, train.crim_median, k = 100)
table(pred.knn, crim_median.test)
```

```
##          crim_median.test
## pred.knn  0
##           0 49
##           1  2
```

```
mean(pred.knn == crim_median.test)
```

```
## [1] 0.9607843
```

When using forward selection params all classified incorrectly for all k.