

Stat 897 Fall 2017 Data Analysis Assignment 9

Penn State

Due November 12, 2017

In this assignment we will use the OJ data found in the ISLR library.

```
library(ISLR)
library(e1071)
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
library(knitr)
```

```
data("OJ")
str(OJ)
```

```
## 'data.frame': 1070 obs. of 18 variables:
## $ Purchase : Factor w/ 2 levels "CH","MM": 1 1 1 2 1 1 1 1 1 1 ...
## $ WeekofPurchase: num 237 239 245 227 228 230 232 234 235 238 ...
## $ StoreID : num 1 1 1 1 7 7 7 7 7 7 ...
## $ PriceCH : num 1.75 1.75 1.86 1.69 1.69 1.69 1.69 1.75 1.75 1.75 ...
## $ PriceMM : num 1.99 1.99 2.09 1.69 1.69 1.99 1.99 1.99 1.99 1.99 ...
## $ DiscCH : num 0 0 0.17 0 0 0 0 0 0 0 ...
## $ DiscMM : num 0 0.3 0 0 0 0 0.4 0.4 0.4 0.4 ...
## $ SpecialCH : num 0 0 0 0 0 0 1 1 0 0 ...
## $ SpecialMM : num 0 1 0 0 0 1 1 0 0 0 ...
## $ LoyalCH : num 0.5 0.6 0.68 0.4 0.957 ...
## $ SalePriceMM : num 1.99 1.69 2.09 1.69 1.69 1.99 1.59 1.59 1.59 1.59 ...
## $ SalePriceCH : num 1.75 1.75 1.69 1.69 1.69 1.69 1.69 1.75 1.75 1.75 ...
## $ PriceDiff : num 0.24 -0.06 0.4 0 0 0.3 -0.1 -0.16 -0.16 -0.16 ...
## $ Store7 : Factor w/ 2 levels "No","Yes": 1 1 1 1 2 2 2 2 2 2 ...
## $ PctDiscMM : num 0 0.151 0 0 0 ...
## $ PctDiscCH : num 0 0 0.0914 0 0 ...
## $ ListPriceDiff : num 0.24 0.24 0.23 0 0 0.3 0.3 0.24 0.24 0.24 ...
## $ STORE : num 1 1 1 1 0 0 0 0 0 0 ...
```

We see that Purchase has 2 levels and we proceed to apply SVM on the data set.

(a) Create a training set containing a random sample of 800 observations, and a test set containing the remaining observations.

```
set.seed(35)
train=sample(nrow(OJ),800)
OJ.train = OJ[train,]
OJ.test = OJ[-train,]
```

(b) Fit a support vector classifier to the training data using `cost=0.01`, with `Purchase` as the response and the other variables as predictors. Use the `summary()` function to produce summary statistics, and describe the results obtained.

A support vector classifier would mean that we run `svm` with `kernel=linear`

```
svm_1 = svm(Purchase~.,data=OJ.train,kernel="linear",cost=0.01)
summary(svm_1)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = OJ.train, kernel = "linear",
##      cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:  0.01
##   gamma:  0.05555556
##
## Number of Support Vectors:  426
##
## ( 214 212 )
##
##
## Number of Classes:  2
##
## Levels:
##   CH MM
```

The summary tells us that the model selects 426 out of 800 observations as support points. The two levels contribute 214 and 212 data points.

(c) What are the training and test error rates?

```
# Training Error
pred.train = predict(svm_1, newdata=OJ.train)
table(pred=pred.train, truth=OJ.train$Purchase)
```

```
##      truth
## pred  CH  MM
##   CH 433  73
##   MM  59 235
```

```
cm=confusionMatrix(data = pred.train, reference = OJ.train$Purchase)
cm$byClass
```

##	Sensitivity	Specificity	Pos Pred Value
##	0.8800813	0.7629870	0.8557312
##	Neg Pred Value	Precision	Recall
##	0.7993197	0.8557312	0.8800813
##	F1	Prevalence	Detection Rate
##	0.8677355	0.6150000	0.5412500

```
## Detection Prevalence    Balanced Accuracy
##           0.6325000           0.8215342
mean(pred.train == OJ.train$Purchase)

## [1] 0.835
# Test Error
pred.test = predict(svm_1,newdata=OJ.test)
table(pred=pred.test, truth=OJ.test$Purchase)

##      truth
## pred CH  MM
##   CH 147  37
##   MM  14  72

cm=confusionMatrix(data = pred.test, reference = OJ.test$Purchase)
cm$byClass

##      Sensitivity      Specificity      Pos Pred Value
##      0.9130435      0.6605505      0.7989130
##      Neg Pred Value      Precision      Recall
##      0.8372093      0.7989130      0.9130435
##      F1      Prevalence      Detection Rate
##      0.8521739      0.5962963      0.5444444
## Detection Prevalence    Balanced Accuracy
##           0.6814815           0.7867970
mean(pred.test == OJ.test$Purchase)
```

```
## [1] 0.8111111
```

In summary we see the following results for linear kernel:

	Training	Test
Mean	0.835	0.8111111
Sensitivity	0.8800813	0.9130435
Specificity	0.7629870	0.6605505
Precision	0.8557312	0.7989130

We see that while the sensitivity is slightly higher, the overall accuracy drops for the test data as compared to the training data.

(d) Use the `tune()` function to select an optimal cost. Consider values in the range 0.01 to 10.

```
set.seed(35)
cost_range = c(.01, .02, .03, .05, .2, .4, .7, 1,2,3,4,5,6,7,8,9,10)

tune.out = tune(svm, Purchase~., data=OJ.train, kernel="linear", ranges=list(cost=cost_range))
summary(tune.out)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
```

```
## 0.03
##
## - best performance: 0.16
##
## - Detailed performance results:
##      cost    error dispersion
## 1  0.01 0.16750 0.03641962
## 2  0.02 0.16625 0.03438447
## 3  0.03 0.16000 0.03106892
## 4  0.05 0.16375 0.03653860
## 5  0.20 0.16375 0.03653860
## 6  0.40 0.16375 0.03653860
## 7  0.70 0.16375 0.03972562
## 8  1.00 0.16500 0.03944053
## 9  2.00 0.16875 0.04093101
## 10 3.00 0.16875 0.03784563
## 11 4.00 0.16750 0.03593976
## 12 5.00 0.17000 0.04090979
## 13 6.00 0.17000 0.04090979
## 14 7.00 0.16875 0.03919768
## 15 8.00 0.16875 0.03919768
## 16 9.00 0.16875 0.03919768
## 17 10.00 0.16875 0.03919768
```

From the results we can say that the optimal results are achieved when cost=0.03. We get error: 0.16

(e) Compute the training and test error rates using this new value for cost.

```
bestmod =tune.out$best.model
summary (bestmod )
```

```
##
## Call:
## best.tune(method = svm, train.x = Purchase ~ ., data = OJ.train,
##      ranges = list(cost = cost_range), kernel = "linear")
##
##
## Parameters:
##      SVM-Type:  C-classification
##      SVM-Kernel:  linear
##              cost:  0.03
##              gamma:  0.05555556
##
## Number of Support Vectors:  368
##
## ( 185 183 )
##
##
## Number of Classes:  2
##
## Levels:
##      CH MM
```

Training Error

```
pred.train = predict(bestmod, newdata=OJ.train)
table(pred=pred.train, truth=OJ.train$Purchase)
```

```
##      truth
## pred CH  MM
##   CH 434  70
##   MM  58 238
```

```
cm=confusionMatrix(data = pred.train, reference = OJ.train$Purchase)
cm$byClass
```

```
##      Sensitivity      Specificity      Pos Pred Value
##      0.8821138      0.7727273      0.8611111
##      Neg Pred Value      Precision      Recall
##      0.8040541      0.8611111      0.8821138
##      F1      Prevalence      Detection Rate
##      0.8714859      0.6150000      0.5425000
## Detection Prevalence      Balanced Accuracy
##      0.6300000      0.8274205
```

```
mean(pred.train == OJ.train$Purchase)
```

```
## [1] 0.84
```

Test Error

```
pred.test = predict(bestmod, newdata=OJ.test)
table(pred=pred.test, truth=OJ.test$Purchase)
```

```
##      truth
## pred CH  MM
##   CH 146  36
##   MM  15  73
```

```
cm=confusionMatrix(data = pred.test, reference = OJ.test$Purchase)
cm$byClass
```

```
##      Sensitivity      Specificity      Pos Pred Value
##      0.9068323      0.6697248      0.8021978
##      Neg Pred Value      Precision      Recall
##      0.8295455      0.8021978      0.9068323
##      F1      Prevalence      Detection Rate
##      0.8513120      0.5962963      0.5407407
## Detection Prevalence      Balanced Accuracy
##      0.6740741      0.7882785
```

```
mean(pred.test == OJ.test$Purchase)
```

```
## [1] 0.8111111
```

In summary we see the following results for bestmodel that uses linear kernel:

	Training	Test
--	----------	------

Mean	0.84	0.8111111
Sensitivity	0.8821138	0.9068323
Specificity	0.7727273	0.6697248
Precision	0.8611111	0.8021978

(f) Repeat parts (b) through (e) using a support vector machine with a radial kernel. Use the default value for gamma.

```
# b. SVM with radial kernel and cost=0.01
svm_r = svm(Purchase~.,data=OJ.train,kernel="radial",cost=0.01)
summary(svm_r)

##
## Call:
## svm(formula = Purchase ~ ., data = OJ.train, kernel = "radial",
##     cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##     cost:  0.01
##    gamma:  0.05555556
##
## Number of Support Vectors:  617
##
##   ( 309 308 )
##
##
## Number of Classes:  2
##
## Levels:
##   CH MM
```

The summary tells us that the model selects 617 out of 800 observations as support points. The two levels contribute 309 and 308 data points.

```
# c. Training and test error rates
# Training Error
pred.train = predict(svm_r, newdata=OJ.train)
table(pred=pred.train, truth=OJ.train$Purchase)

##      truth
## pred  CH  MM
##   CH 492 308
##   MM   0   0

cm=confusionMatrix(data = pred.train, reference = OJ.train$Purchase)
cm$byClass

##      Sensitivity      Specificity      Pos Pred Value
##      1.0000000      0.0000000      0.6150000
##      Neg Pred Value      Precision      Recall
##      NaN      0.6150000      1.0000000
##      F1      Prevalence      Detection Rate
##      0.7616099      0.6150000      0.6150000
## Detection Prevalence      Balanced Accuracy
##      1.0000000      0.5000000

mean(pred.train == OJ.train$Purchase)
```

```
## [1] 0.615
```

```
# Test Error
```

```
pred.test = predict(svm_r,newdata=OJ.test)
table(pred=pred.test, truth=OJ.test$Purchase)
```

```
##      truth
## pred  CH  MM
##   CH 161 109
##   MM   0   0
```

```
cm=confusionMatrix(data = pred.test, reference = OJ.test$Purchase)
cm$byClass
```

```
##          Sensitivity      Specificity      Pos Pred Value
##          1.0000000      0.0000000      0.5962963
##      Neg Pred Value      Precision      Recall
##          NaN      0.5962963      1.0000000
##          F1      Prevalence      Detection Rate
##          0.7470998      0.5962963      0.5962963
## Detection Prevalence      Balanced Accuracy
##          1.0000000      0.5000000
```

```
mean(pred.test == OJ.test$Purchase)
```

```
## [1] 0.5962963
```

In summary we see the following results for radial kernel:

	Training	Test
Mean	0.615	0.5962963
Sensitivity	1.0000000	1.0000000
Specificity	0.0000000	0.0000000
Precision	0.6150000	0.5962963

The results have become worse.

```
# d. tune on radial
```

```
cost_range = c(.01, .02, .03, .05, .2, .4, .7, 1,2,3,4,5,6,7,8,9,10)
set.seed(35)
tune.out = tune(svm, Purchase~., data=OJ.train, kernel="radial", ranges=list(cost=cost_range))
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     1
##
## - best performance: 0.1625
##
## - Detailed performance results:
##   cost  error dispersion
## 1  0.01 0.38500 0.04632314
## 2  0.02 0.38500 0.04632314
## 3  0.03 0.36750 0.05957022
## 4  0.05 0.19875 0.02791978
```

```
## 5 0.20 0.17500 0.03632416
## 6 0.40 0.17125 0.03283481
## 7 0.70 0.16500 0.03525699
## 8 1.00 0.16250 0.03486083
## 9 2.00 0.16375 0.03356689
## 10 3.00 0.17000 0.03917553
## 11 4.00 0.17250 0.03670453
## 12 5.00 0.17125 0.03537988
## 13 6.00 0.17250 0.03525699
## 14 7.00 0.17375 0.03304563
## 15 8.00 0.17500 0.03227486
## 16 9.00 0.17750 0.03162278
## 17 10.00 0.17750 0.03162278
```

From the results we can say that the optimal results are achieved when cost=1. We get error: 0.1625

```
# e. training and test error rates using best model
```

```
bestmod =tune.out$best.model
```

```
summary (bestmod )
```

```
##
```

```
## Call:
```

```
## best.tune(method = svm, train.x = Purchase ~ ., data = OJ.train,
```

```
## ranges = list(cost = cost_range), kernel = "radial")
```

```
##
```

```
##
```

```
## Parameters:
```

```
## SVM-Type: C-classification
```

```
## SVM-Kernel: radial
```

```
## cost: 1
```

```
## gamma: 0.05555556
```

```
##
```

```
## Number of Support Vectors: 360
```

```
##
```

```
## ( 183 177 )
```

```
##
```

```
##
```

```
## Number of Classes: 2
```

```
##
```

```
## Levels:
```

```
## CH MM
```

```
# Training Error
```

```
pred.train = predict(bestmod, newdata=OJ.train)
```

```
table(pred=pred.train, truth=OJ.train$Purchase)
```

```
## truth
```

```
## pred CH MM
```

```
## CH 451 73
```

```
## MM 41 235
```

```
cm=confusionMatrix(data = pred.train, reference = OJ.train$Purchase)
```

```
cm$byClass
```

```
##          Sensitivity          Specificity          Pos Pred Value
##          0.9166667          0.7629870          0.8606870
##          Neg Pred Value          Precision          Recall
```



```
##           0.8514493           0.8606870           0.9166667
##           F1           Prevalence           Detection Rate
##           0.8877953           0.6150000           0.5637500
## Detection Prevalence      Balanced Accuracy
##           0.6550000           0.8398268
```

```
mean(pred.train == OJ.train$Purchase)
```

```
## [1] 0.8575
```

```
# Test Error
```

```
pred.test = predict(bestmod,newdata=OJ.test)
table(pred=pred.test, truth=OJ.test$Purchase)
```

```
##      truth
## pred CH  MM
##   CH 147  38
##   MM  14  71
```

```
cm=confusionMatrix(data = pred.test, reference = OJ.test$Purchase)
cm$byClass
```

```
##           Sensitivity           Specificity           Pos Pred Value
##           0.9130435           0.6513761           0.7945946
##           Neg Pred Value           Precision           Recall
##           0.8352941           0.7945946           0.9130435
##           F1           Prevalence           Detection Rate
##           0.8497110           0.5962963           0.5444444
## Detection Prevalence      Balanced Accuracy
##           0.6851852           0.7822098
```

```
mean(pred.test == OJ.test$Purchase)
```

```
## [1] 0.8074074
```

In summary we see the following results for bestmodel that uses radial kernel:

	Training	Test
Mean	0.8575	0.8074074
Sensitivity	0.9166667	0.9130435
Specificity	0.7629870	0.6513761
Precision	0.8606870	0.7945946

(g) Repeat parts (b) through (e) using a support vector machine with a polynomial kernel. Set degree=2.

```
# b. SVM with polynomial kernel and cost=0.01
svm_p = svm(Purchase~.,data=OJ.train,kernel="polynomial", degree=2,cost=0.01)
summary(svm_p)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = OJ.train, kernel = "polynomial",
##      degree = 2, cost = 0.01)
##
##
## Parameters:
##   SVM-Type:  C-classification
```

```
## SVM-Kernel: polynomial
##      cost: 0.01
##      degree: 2
##      gamma: 0.05555556
##      coef.0: 0
##
## Number of Support Vectors: 620
##
## ( 312 308 )
##
##
## Number of Classes: 2
##
## Levels:
## CH MM
```

The summary tells us that the model selects 620 out of 800 observations as support points. The two levels contribute 312 and 308 data points.

```
# c. Training and test error rates
# Training Error
pred.train = predict(svm_p, newdata=OJ.train)
table(pred=pred.train, truth=OJ.train$Purchase)
```

```
##      truth
## pred CH MM
## CH 491 295
## MM   1  13
```

```
cm=confusionMatrix(data = pred.train, reference = OJ.train$Purchase)
cm$byClass
```

##	Sensitivity	Specificity	Pos Pred Value
##	0.99796748	0.04220779	0.62468193
##	Neg Pred Value	Precision	Recall
##	0.92857143	0.62468193	0.99796748
##	F1	Prevalence	Detection Rate
##	0.76838811	0.61500000	0.61375000
##	Detection Prevalence	Balanced Accuracy	
##	0.98250000	0.52008764	

```
mean(pred.train == OJ.train$Purchase)
```

```
## [1] 0.63
```

```
# Test Error
pred.test = predict(svm_p,newdata=OJ.test)
table(pred=pred.test, truth=OJ.test$Purchase)
```

```
##      truth
## pred CH MM
## CH 161 109
## MM   0   0
```

```
cm=confusionMatrix(data = pred.test, reference = OJ.test$Purchase)
cm$byClass
```

##	Sensitivity	Specificity	Pos Pred Value
----	-------------	-------------	----------------

```
##          1.0000000          0.0000000          0.5962963
##      Neg Pred Value          Precision          Recall
##          NaN          0.5962963          1.0000000
##          F1          Prevalence          Detection Rate
##      0.7470998          0.5962963          0.5962963
## Detection Prevalence    Balanced Accuracy
##          1.0000000          0.5000000
```

```
mean(pred.test == OJ.test$Purchase)
```

```
## [1] 0.5962963
```

In summary we see the following results for polynomial kernel:

	Training	Test
Mean	0.63	0.5962963
Sensitivity	0.99796748	1.0000000
Specificity	0.04220779	0.0000000
Precision	0.62468193	0.5962963

```
# d. tune on polynomial
cost_range = c(.01, .02, .03, .05, .2, .4, .7, 1,2,3,4,5,6,7,8,9,10)
set.seed(35)
tune.out = tune(svm, Purchase~., data=OJ.train, kernel="polynomial", degree=2, ranges=list(cost=cost_range))
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     9
##
## - best performance: 0.16875
##
## - Detailed performance results:
##   cost  error dispersion
## 1  0.01 0.38500 0.04632314
## 2  0.02 0.35875 0.04788949
## 3  0.03 0.35625 0.04938862
## 4  0.05 0.33500 0.04958158
## 5  0.20 0.24250 0.04571956
## 6  0.40 0.20250 0.02934469
## 7  0.70 0.20000 0.03173239
## 8  1.00 0.19250 0.02958040
## 9  2.00 0.18750 0.03908680
## 10 3.00 0.18125 0.03596391
## 11 4.00 0.17750 0.03944053
## 12 5.00 0.17125 0.03634805
## 13 6.00 0.17750 0.04241004
## 14 7.00 0.17875 0.03682259
## 15 8.00 0.16875 0.03875224
## 16 9.00 0.16875 0.03963812
## 17 10.00 0.17125 0.03866254
```

From the results we can say that the optimal results are achieved when cost=9. We get error: 0.16875

```

# e. training and test error rates using best model
bestmod =tune.out$best.model
summary (bestmod )

##
## Call:
## best.tune(method = svm, train.x = Purchase ~ ., data = OJ.train,
##   ranges = list(cost = cost_range), kernel = "polynomial",
##   degree = 2)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##     cost:  9
##   degree:  2
##   gamma:  0.05555556
##   coef.0:  0
##
## Number of Support Vectors:  338
##
## ( 173 165 )
##
##
## Number of Classes:  2
##
## Levels:
##  CH MM
# Training Error
pred.train = predict(bestmod, newdata=OJ.train)
table(pred=pred.train, truth=OJ.train$Purchase)

##      truth
## pred  CH  MM
##   CH 450  74
##   MM  42 234

cm=confusionMatrix(data = pred.train, reference = OJ.train$Purchase)
cm$byClass

##      Sensitivity      Specificity      Pos Pred Value
##      0.9146341      0.7597403      0.8587786
##      Neg Pred Value      Precision      Recall
##      0.8478261      0.8587786      0.9146341
##      F1      Prevalence      Detection Rate
##      0.8858268      0.6150000      0.5625000
## Detection Prevalence      Balanced Accuracy
##      0.6550000      0.8371872

mean(pred.train == OJ.train$Purchase)

## [1] 0.855
# Test Error
pred.test = predict(bestmod,newdata=OJ.test)

```

```
table(pred=pred.test, truth=OJ.test$Purchase)
```

```
##      truth
## pred CH  MM
##   CH 151  43
##   MM  10  66
```

```
cm=confusionMatrix(data = pred.test, reference = OJ.test$Purchase)
cm$byClass
```

```
##          Sensitivity          Specificity          Pos Pred Value
##          0.9378882          0.6055046          0.7783505
##          Neg Pred Value          Precision          Recall
##          0.8684211          0.7783505          0.9378882
##          F1          Prevalence          Detection Rate
##          0.8507042          0.5962963          0.5592593
## Detection Prevalence          Balanced Accuracy
##          0.7185185          0.7716964
```

```
mean(pred.test == OJ.test$Purchase)
```

```
## [1] 0.8037037
```

In summary we see the following results for bestmodel that uses radial kernel:

	Training	Test
Mean	0.855	0.8037037
Sensitivity	0.9146341	0.9378882
Specificity	0.7597403	0.6055046
Precision	0.8587786	0.7783505

(h) Overall, which approach seems to give the best results on this data?

We will compare the test data metrics for the best models retrieved from linear, radial and polynomial kernels.

	Test (poly, deg:2)	Test Radial	Test Linear
Mean	0.8037037	0.8074074	0.8111111
Sensitivity	0.9378882	0.9130435	0.9068323
Specificity	0.6055046	0.6513761	0.6697248
Precision	0.7783505	0.7945946	0.8021978
CV Error	0.16875	0.1625	0.16

We get the best mean/precision and specificity with linear kernel. The sensitivity is highest for the polynomial (degree=2) kernel. The CV error is only slightly better for the linear model.